

# Assignment 1

## Title: Exploring the MovieLens 1M Dataset

Name: Pranjal Rane NUID: 002756852

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

### Importing Data and Exploratory Analysis

In [2]:

```
unames = ['user_id', 'gender', 'age', 'occupation', 'zip']
users = pd.read_table('ml-1m/users.dat', sep='::', header=None, names=unames, engine='python', encoding='latin-1')

rnames = ['user_id', 'movie_id', 'rating', 'timestamp']
ratings = pd.read_table('ml-1m/ratings.dat', sep='::', header=None, names=rnames, engine='python', encoding='latin-1')

mnames = ['movie_id', 'title', 'genres']
movies = pd.read_table('ml-1m/movies.dat', sep='::', header=None, names=mnames, engine='python', encoding='latin-1')
```

In [3]:

```
users[:5]
```

Out[3]:

	user_id	gender	age	occupation	zip
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

In [4]:

```
ratings[:5]
```

Out[4]:

	user_id	movie_id	rating	timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

In [5]:

```
movies[:5]
```

Out[5]:

	movie_id	title	genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

## Merging the three tables

In [6]:

```
data = pd.merge(pd.merge(ratings, users), movies)
```

In [7]:

```
data.head(1)
```

Out[7]:

	user_id	movie_id	rating	timestamp	gender	age	occupation	zip	title	genres
0	1	1193	5	978300760	F	1	10 48067		One Flew Over the Cuckoo's Nest (1975)	Drama

## Objective 1

An aggregate of movie ratings by men of age above 25 for each particular genre

In [8]:

```
filtered_data = data[(data['gender'] == 'M') & (data['age'] > 25)]
filtered_data = filtered_data.assign(genres=filtered_data['genres'].str.split('|')).explode('genres')
genre_ratings = filtered_data.groupby('genres')['rating'].mean().reset_index()
genre_ratings_sorted = genre_ratings.sort_values(by='rating', ascending=False)

genre_ratings_sorted
```

Out[8]:

	genres	rating
9	Film-Noir	4.117140
6	Documentary	3.950192
16	War	3.940634
7	Drama	3.812309
5	Crime	3.764249
12	Mystery	3.759347
2	Animation	3.721569
17	Western	3.708494
11	Musical	3.700242
13	Romance	3.659748
15	Thriller	3.644025
4	Comedy	3.565456

0	genres	3.56184
1	Adventure	3.538637
14	Sci-Fi	3.509693
8	Fantasy	3.490408
3	Children's	3.475314
10	Horror	3.241089

## Objective 2

The top 5 ranked movies by the most number of ratings (not the highest rating)

In [9]:

```
most Rated = data.groupby('title').size().sort_values(ascending=False)[:5]
most Rated
```

Out[9]:

```
title
American Beauty (1999)          3428
Star Wars: Episode IV - A New Hope (1977)    2991
Star Wars: Episode V - The Empire Strikes Back (1980)  2990
Star Wars: Episode VI - Return of the Jedi (1983)    2883
Jurassic Park (1993)            2672
dtype: int64
```

## Objective 3

Average movie ratings between users of different age groups (<18, 18-30, 30-50, 50-70, 70>)

In [10]:

```
bins = [0, 18, 30, 50, 70, float('inf')]
labels = ['<18', '18-30', '30-50', '50-70', '70>']
data['age_group'] = pd.cut(data['age'], bins=bins, labels=labels, right=False)
age_group_ratings = data.groupby('age_group')['rating'].mean().reset_index()

age_group_ratings
```

Out[10]:

	age_group	rating
0	<18	3.549520
1	18-30	3.533299
2	30-50	3.624050
3	50-70	3.732677
4	70>	NaN

## Objective 4

Pick a movie of your choice and for all movies of the same year, provide a breakdown of the number of unique movies rated by 3 ranges of age of reviewers (a) under 18 (b) 19 to 45 (c) Above 45.

In [11]:

```
chosen_movie = "Dumb & Dumber (1994)"

data['release_year'] = data['title'].str.extract(r'\((\d{4})\)').astype(int)

chosen_movie_year = data[data['title'] == chosen_movie]['release_year'].values[0]
```

```

same_year_movies = data[data['release_year'] == chosen_movie_year]

bins = [0, 18, 45, float('inf')]
labels = ['Under 18', '19 to 45', 'Above 45']
same_year_movies['age_group'] = pd.cut(same_year_movies['age'], bins=bins, labels=labels,
, right=False)

age_group_movie_counts = same_year_movies.groupby('age_group')['movie_id'].nunique().reset_index(name='unique_movies_count')

age_group_movie_counts

```

Out[11]:

	age_group	unique_movies_count
0	Under 18	154
1	19 to 45	241
2	Above 45	226

## Objective 5

A function that takes in a `user_id` and a `movie_id`, and returns a list of all the other movies that the user rated similarly to the given movie, i.e. with the same rating. Demonstrate that your function works.

In [12]:

```

def find_similarly_rated_movies(user_id, movie_id, data):
    user_rating = data[(data['user_id'] == user_id) & (data['movie_id'] == movie_id)]['rating'].values[0]

    user_data = data[data['user_id'] == user_id]

    similarly_rated_movies = user_data[user_data['rating'] == user_rating]['movie_id'].unique()

    similarly_rated_movies = similarly_rated_movies[similarly_rated_movies != movie_id]

    return similarly_rated_movies

user_id_example = 12
movie_id_example = 1193

similar_movies = find_similarly_rated_movies(user_id_example, movie_id_example, data)
print(f"User {user_id_example} rated movies {similar_movies} similarly to movie {movie_id_example}.")

selected = movies[movies["movie_id"].isin(similar_movies)][['movie_id', 'title']].value_counts().reset_index()[['movie_id', 'title']]
selected

```

User 12 rated movies [3265 3897 3658 1303 999] similarly to movie 1193.

Out[12]:

	movie_id	title
0	999	2 Days in the Valley (1996)
1	1303	Man Who Would Be King, The (1975)
2	3265	Hard-Boiled (Lashou shentan) (1992)
3	3658	Quatermass and the Pit (1967)
4	3897	Almost Famous (2000)


## Objective 6

Some other statistic, figure, aggregate, or plot that you created using this dataset, along with a short description of what interesting observations you derived from it.

In [13]:

```
average_ratings = data.groupby('movie_id')['rating'].mean()

plt.figure(figsize=(10, 6))
plt.hist(average_ratings, bins=20, color='skyblue', edgecolor='black')
plt.title('Distribution of Average Movie Ratings')
plt.xlabel('Average Rating')
plt.ylabel('Number of Movies')
plt.grid(axis='y', alpha=0.75)
plt.show()
```



From the above histogram we can infer that:

1. A peak at a certain rating range (3.0-3.5). This peak indicates the most common average rating that movies receive.
2. The distribution is slightly right-skewed, indicating ratings above 2.5 are more common.

In [14]:


```
average_ratings = data.groupby('movie_id')['rating'].mean()

ratings_count = data.groupby('movie_id')['rating'].count()

movie_stats = pd.DataFrame({
    'average_rating': average_ratings,
    'number_of_ratings': ratings_count
})

plt.figure(figsize=(10, 6))
plt.scatter(movie_stats['number_of_ratings'], movie_stats['average_rating'], alpha=0.5, color='blue')
plt.title('Average Rating vs. Number of Ratings')
plt.xlabel('Number of Ratings')
```

```
plt.ylabel('Average Rating')
plt.grid(True)
plt.show()
```



From the above scatter plot we can infer that:

1. Areas with more points (3.0 - 4.0) indicate a common number of ratings and average rating scores
2. Movies with ratings on extreme end of the spectrum has considerably less number of ratings than those with ratings which are part of the bell curve